

Created on Sat Nov 06 16:00:00 2021. @author: Cedric Yu

Climate Change: Global Temperature

We study the global temperatures dataset available on:

<https://www.kaggle.com/berkeleyearth/climate-change-earth-surface-temperature-data>

We focus on the time series 'LandAverageTemperature', which is the monthly global average land temperature in Celsius.

We first explore the given relative absolute uncertainty of the time series, and find that it is large during the early record 1750-1849. Because of this, we decide to take the yearly averages and only use the data from 1850-2015.

By eyeballing and using the augmented Dickey-Fuller test, we see a clear increasing trend (hence global 'warming'), which is rendered stationary by taking the first difference.

Next, we study the auto-correlation function (ACF) and partial auto-correlation function (PACF), we find that the first difference can be described by an AR(3) model. We fit AR(3) models on the series, using fixed partitioning and rolling forecast, and compare the model performance.

Preamble

```
In [1]: import pandas as pd
# Make the output look better
pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
# pd.set_option('display.width', 1000)
pd.options.mode.chained_assignment = None # default='warn' # ignores warning about dropping columns inplace

import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt

import os
os.chdir(r'C:\Users\Cedric Yu\Desktop\Works\13_time_series_global_temperatures')
```

load dataset and parse datetime

```
In [2]: """# import GlobalTemperatures dataset and parse datetime column 'dt'"""

from datetime import datetime
def parser(x):
    return datetime.strptime(x, '%Y-%m-%d')

global_temperatures = pd.read_csv('GlobalTemperatures.csv', header = [0], parse_dates = [0], date_parser=parser)
```

Preliminary observations

```
In [4]: global_temperatures.tail()
```

```
Out[4]:
```

	dt	LandAverageTemperature	LandAverageTemperatureUncertainty	LandMaxTemperature	LandMaxTemperatureUncertainty	LandMinTemperature	LandMinTemperatureUncertainty
3187	2015-08-01	14.755	0.072	20.699	0.110	8.811	0.088
3188	2015-09-01	12.999	0.079	18.845	0.088	6.153	0.093
3189	2015-10-01	10.801	0.102	16.450	0.059	5.152	0.093
3190	2015-11-01	7.433	0.119	12.892	0.093	2.974	0.088

	dt	LandAverageTemperature	LandAverageTemperatureUncertainty	LandMaxTemperature	LandMaxTemperatureUncertainty	LandMinTemperature	LandMinTemperatureUncertainty
3191	2015-12-01	5.518	0.100	10.725	0.154		

```
In [5]: """# dtype of column 'dt' is datetime64[ns]"""
global_temperatures['dt']
```

```
Out[5]: 0      1750-01-01
1      1750-02-01
2      1750-03-01
3      1750-04-01
4      1750-05-01
...
3187   2015-08-01
3188   2015-09-01
3189   2015-10-01
3190   2015-11-01
3191   2015-12-01
Name: dt, Length: 3192, dtype: datetime64[ns]
```

```
In [6]: """# missing values"""
global_temperatures.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3192 entries, 0 to 3191
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   dt                                    3192 non-null   datetime64[ns]
1   LandAverageTemperature                3180 non-null   float64
2   LandAverageTemperatureUncertainty     3180 non-null   float64
3   LandMaxTemperature                   1992 non-null   float64
4   LandMaxTemperatureUncertainty         1992 non-null   float64
5   LandMinTemperature                   1992 non-null   float64
6   LandMinTemperatureUncertainty         1992 non-null   float64
7   LandAndOceanAverageTemperature       1992 non-null   float64
8   LandAndOceanAverageTemperatureUncertainty 1992 non-null   float64
dtypes: datetime64[ns](1), float64(8)
memory usage: 224.6 KB
```

```
In [7]: """# 'LandAverageTemperature' and 'LandAverageTemperatureUnvertainty' have no NaN after 1752-12-31"""
global_temperatures[global_temperatures[['LandAverageTemperature']].isna().any(axis=1)]['dt']
```

```
Out[7]: 10     1750-11-01
16     1751-05-01
18     1751-07-01
21     1751-10-01
22     1751-11-01
23     1751-12-01
25     1752-02-01
28     1752-05-01
29     1752-06-01
30     1752-07-01
31     1752-08-01
32     1752-09-01
Name: dt, dtype: datetime64[ns]
```

```
In [8]: # 'LandAverageTemperature' and 'LandAverageTemperatureUnvertainty' have no NaN after 1752-12-31
global_temperatures[global_temperatures['dt'] > pd.Timestamp(1752,12,31)].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3156 entries, 36 to 3191
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   dt                                    3156 non-null   datetime64[ns]
1   LandAverageTemperature                3156 non-null   float64
2   LandAverageTemperatureUncertainty     3156 non-null   float64
3   LandMaxTemperature                   1992 non-null   float64
4   LandMaxTemperatureUncertainty         1992 non-null   float64
5   LandMinTemperature                   1992 non-null   float64
6   LandMinTemperatureUncertainty         1992 non-null   float64
```

```

7   LandAndOceanAverageTemperature      1992 non-null   float64
8   LandAndOceanAverageTemperatureUncertainty  1992 non-null   float64
dtypes: datetime64[ns](1), float64(8)
memory usage: 246.6 KB

```

```

In [9]: """# No NaN at all after 1849-12-31"""
global_temperatures[global_temperatures[['LandMaxTemperature']].isna().any(axis=1)]['dt']

```

```

Out[9]: 0      1750-01-01
1      1750-02-01
2      1750-03-01
3      1750-04-01
4      1750-05-01
...
1195   1849-08-01
1196   1849-09-01
1197   1849-10-01
1198   1849-11-01
1199   1849-12-01
Name: dt, Length: 1200, dtype: datetime64[ns]

```

```

In [10]: global_temperatures[global_temperatures['dt'] > pd.Timestamp(1849,12,31)].info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1992 entries, 1200 to 3191
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   dt                                    1992 non-null   datetime64[ns]
1   LandAverageTemperature                1992 non-null   float64
2   LandAverageTemperatureUncertainty     1992 non-null   float64
3   LandMaxTemperature                   1992 non-null   float64
4   LandMaxTemperatureUncertainty         1992 non-null   float64
5   LandMinTemperature                   1992 non-null   float64
6   LandMinTemperatureUncertainty         1992 non-null   float64
7   LandAndOceanAverageTemperature        1992 non-null   float64
8   LandAndOceanAverageTemperatureUncertainty  1992 non-null   float64
dtypes: datetime64[ns](1), float64(8)
memory usage: 155.6 KB

```

Looking at the uncertainties

LandAverageTemperature: plots by month

```

In [3]: temperatures_1753_avg = global_temperatures[global_temperatures['dt'] > pd.Timestamp(1752,12,31)][['dt', 'LandAverageTemperature']]
temperatures_1850_all = global_temperatures[global_temperatures['dt'] > pd.Timestamp(1849,12,31)]


```

```

In [15]: """# plot LandAverageTemperature (green) and error bar 'LandAverageTemperatureUncertainty' (light green)"""

plt.figure(dpi=100)
plt.errorbar('dt', 'LandAverageTemperature', yerr='LandAverageTemperatureUncertainty', data=temperatures_1753_avg)
ax1 = plt.gca()
ax1.set_title('Land Average Temperature by Month')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)

```


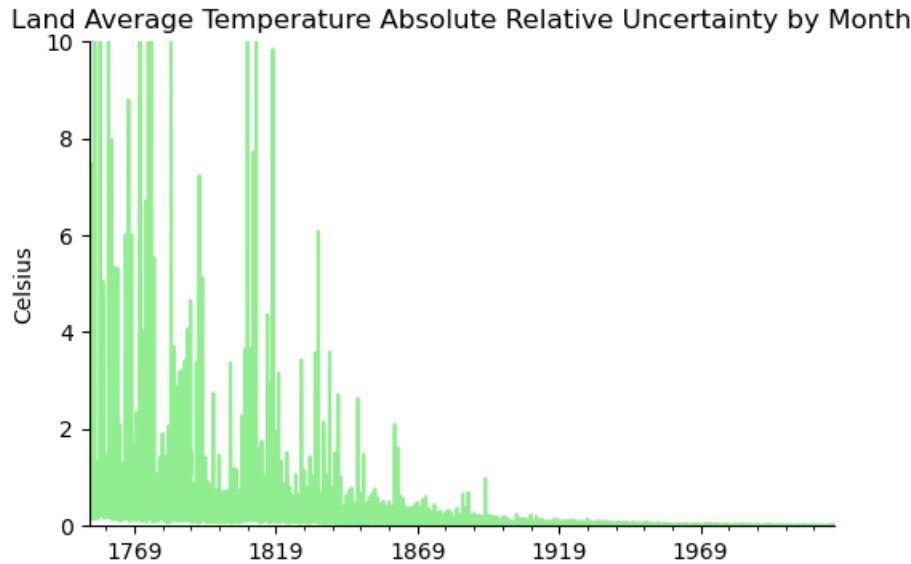


Absolute Relative Uncertainty

```
In [35]: """relative uncertainty"""
temperatures_1753_avg['LandAverageTemperatureRelativeUncertainty'] = (temperatures_1753_avg['LandAverageTemperatu

"""# large relative error in earlier observations"""
plt.figure(dpi=100)
temperatures_1753_avg.plot(x = 'dt', y='LandAverageTemperatureRelativeUncertainty', color='lightgreen', ax = plt.
ax1 = plt.gca()
ax1.set_title('Land Average Temperature Absolute Relative Uncertainty by Month')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.set_ylim(0,10)
ax1.set_xlabel(None)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
ax1.get_legend().remove()
# plt.savefig(r'pLots/LandAverageTemperatureRelativeUncertainty_month1', dpi=150)

plt.figure(dpi=100)
temperatures_1753_avg.plot(x = 'dt', y='LandAverageTemperatureRelativeUncertainty', color='lightgreen', ax = plt.
ax1 = plt.gca()
ax1.set_ylim(0,1)
ax1.set_title('Land Average Temperature Absolute Relative Uncertainty by Month')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.set_xlim(pd.Timestamp(1845, 12, 31))
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
ax1.get_legend().remove()
# plt.savefig(r'pLots/LandAverageTemperatureRelativeUncertainty_month2', dpi=150)
```



LandAverageTemperature: Yearly mean temperature and uncertainty

```
In [4]: def yearly_uncertainty(col):
        return np.sqrt(np.sum(col ** 2)) / len(col)

temperatures_1753_avg_yearly = temperatures_1753_avg[['dt', 'LandAverageTemperature']].groupby(pd.Grouper(key='dt',
temperatures_1753_avg_yearly['LandAverageTemperatureUncertainty'] = temperatures_1753_avg.groupby(pd.Grouper(key=
temperatures_1753_avg_yearly['LandAverageTemperatureRelativeUncertainty'] = (temperatures_1753_avg_yearly['LandAv

temperatures_1753_avg_yearly.head()
```

C:\Users\CEDRIC~1\AppData\Local\Temp\ipykernel_16880\2418994677.py:5: FutureWarning: Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.

```
temperatures_1753_avg_yearly['LandAverageTemperatureUncertainty'] = temperatures_1753_avg.groupby(pd.Grouper(key='dt', freq='1Y'))['dt', 'LandAverageTemperatureUncertainty'].agg(yearly_uncertainty)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\groupby\generic.py:303: FutureWarning: Dropping invalid columns in SeriesGroupBy.agg is deprecated. In a future version, a TypeError will be raised. Before calling .agg, select only columns which should be valid for the aggregating function.

```
results[key] = self.aggregate(func)
```

Out[4]:

	LandAverageTemperature	LandAverageTemperatureUncertainty	LandAverageTemperatureRelativeUncertainty
dt			
1753-12-31	8.388083	0.970134	0.115656
1754-12-31	8.469333	1.091063	0.128825

	LandAverageTemperature	LandAverageTemperatureUncertainty	LandAverageTemperatureRelativeUncertainty
dt			
1755-12-31	8.355583	1.153430	0.138043
1756-12-31	8.849583	0.981928	0.110958
1757-12-31	9.022000	1.180627	0.130861

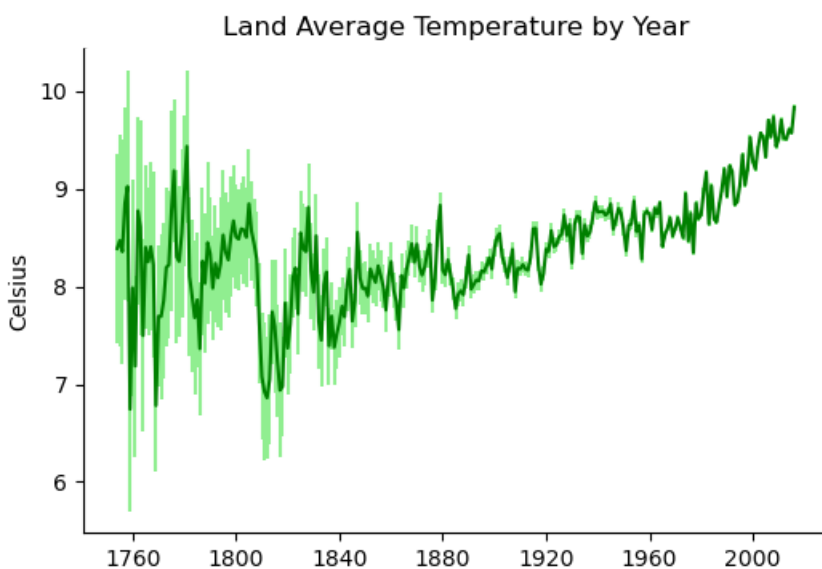
In [27]:

```

"""# plot yearly-averaged LandAverageTemperature and error bar 'LandAverageTemperatureUncertainty'"""
""" relative uncertainty becomes smaller after ca. 1850"""

plt.figure(dpi=100)
plt.errorbar(temperatures_1753_avg_yearly.index, 'LandAverageTemperature', yerr='LandAverageTemperatureUncertainty')
ax1 = plt.gca()
ax1.set_title('Land Average Temperature by Year')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
# ax1.set_ylim(0,10)
ax1.set_xlabel(None)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)

```





In [34]:

```

plt.figure(dpi=100)
temperatures_1753_avg_yearly.plot(y = 'LandAverageTemperatureRelativeUncertainty', color = 'lightgreen', ax = plt)
ax1 = plt.gca()
ax1.set_title('Land Average Temperature Absolute Relative Uncertainty by Year')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
# ax1.set_ylim(0,10)
ax1.set_xlabel(None)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
ax1.get_legend().remove()

plt.figure(dpi=100)
temperatures_1753_avg_yearly.plot(y = 'LandAverageTemperatureRelativeUncertainty', color = 'lightgreen', ax = plt)
ax1 = plt.gca()
ax1.set_title('Land Average Temperature Absolute Relative Uncertainty by Year')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.set_ylim(0,0.04)
ax1.set_xlim(pd.Timestamp(1849, 12, 31))
ax1.set_xlabel(None)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
ax1.get_legend().remove()

```



Land Average/Max/Min Temperature: Yearly mean temperature and uncertainty

```
In [5]: temperatures_1850_all.columns
```

```
Out[5]: Index(['dt', 'LandAverageTemperature', 'LandAverageTemperatureUncertainty',
              'LandMaxTemperature', 'LandMaxTemperatureUncertainty',
              'LandMinTemperature', 'LandMinTemperatureUncertainty',
              'LandAndOceanAverageTemperature',
              'LandAndOceanAverageTemperatureUncertainty'],
              dtype='object')
```

```
In [6]: # annual mean temperature and uncertainty
def yearly_uncertainty(col):
    return np.sqrt(np.sum(col ** 2)) / len(col)
temperatures_1850_all_yearly = temperatures_1850_all[['dt', 'LandAverageTemperature',
              'LandMaxTemperature', 'LandMinTemperature',
              'LandAndOceanAverageTemperature']].groupby(pd.Grouper(key='dt', freq='1Y')).mean()
temperatures_1850_all_yearly['LandAverageTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='dt', freq='1Y')).apply(lambda x: yearly_uncertainty(x['LandAverageTemperature']), axis=1)
temperatures_1850_all_yearly['LandMaxTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='dt', freq='1Y')).apply(lambda x: yearly_uncertainty(x['LandMaxTemperature']), axis=1)
temperatures_1850_all_yearly['LandMinTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='dt', freq='1Y')).apply(lambda x: yearly_uncertainty(x['LandMinTemperature']), axis=1)
temperatures_1850_all_yearly['LandAndOceanAverageTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='dt', freq='1Y')).apply(lambda x: yearly_uncertainty(x['LandAndOceanAverageTemperature']), axis=1)
```

C:\Users\CEDRIC~1\AppData\Local\Temp\ipykernel_16880\177883175.py:7: FutureWarning: Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.
temperatures_1850_all_yearly['LandAverageTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='dt', freq='1Y'))['dt', 'LandAverageTemperatureUncertainty'].agg(yearly_uncertainty)
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\groupby\generic.py:303: FutureWarning: Dropping invalid columns will be deprecated, use dropna=False instead.

```

lums in SeriesGroupBy.agg is deprecated. In a future version, a TypeError will be raised. Before calling .agg, s
elect only columns which should be valid for the aggregating function.
    results[key] = self.aggregate(func)
C:\Users\CEDRIC~1\AppData\Local\Temp\ipykernel_16880\177883175.py:8: FutureWarning: Indexing with multiple keys
(implicitly converted to a tuple of keys) will be deprecated, use a list instead.
    temperatures_1850_all_yearly['LandMaxTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='d
t', freq='1Y'))['dt', 'LandMaxTemperatureUncertainty'].agg(yearly_uncertainty)
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\groupby\generic.py:303: FutureWarning: Dropping invalid co
lums in SeriesGroupBy.agg is deprecated. In a future version, a TypeError will be raised. Before calling .agg, s
elect only columns which should be valid for the aggregating function.
    results[key] = self.aggregate(func)
C:\Users\CEDRIC~1\AppData\Local\Temp\ipykernel_16880\177883175.py:9: FutureWarning: Indexing with multiple keys
(implicitly converted to a tuple of keys) will be deprecated, use a list instead.
    temperatures_1850_all_yearly['LandMinTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Grouper(key='d
t', freq='1Y'))['dt', 'LandMinTemperatureUncertainty'].agg(yearly_uncertainty)
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\groupby\generic.py:303: FutureWarning: Dropping invalid co
lums in SeriesGroupBy.agg is deprecated. In a future version, a TypeError will be raised. Before calling .agg, s
elect only columns which should be valid for the aggregating function.
    results[key] = self.aggregate(func)
C:\Users\CEDRIC~1\AppData\Local\Temp\ipykernel_16880\177883175.py:10: FutureWarning: Indexing with multiple keys
(implicitly converted to a tuple of keys) will be deprecated, use a list instead.
    temperatures_1850_all_yearly['LandAndOceanAverageTemperatureUncertainty'] = temperatures_1850_all.groupby(pd.Gr
ouper(key='dt', freq='1Y'))['dt', 'LandAndOceanAverageTemperatureUncertainty'].agg(yearly_uncertainty)
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\groupby\generic.py:303: FutureWarning: Dropping invalid co
lums in SeriesGroupBy.agg is deprecated. In a future version, a TypeError will be raised. Before calling .agg, s
elect only columns which should be valid for the aggregating function.
    results[key] = self.aggregate(func)

```


In [46]:

```

"""# plot yearly-averaged LandAverageTemperature and error bar 'LandAverageTemperatureUncertainty'"""

plt.figure(dpi=100)
plt.suptitle('Land Temperature Averaged by Year')
plt.errorbar(temperatures_1850_all_yearly.index, 'LandAverageTemperature', yerr='LandAverageTemperatureUncertainty', data
plt.errorbar(temperatures_1850_all_yearly.index, 'LandMaxTemperature', yerr='LandMaxTemperatureUncertainty', data
plt.errorbar(temperatures_1850_all_yearly.index, 'LandMinTemperature', yerr='LandMinTemperatureUncertainty', data
ax1 = plt.gca()
ax1.set_title('Red: Maximum. Green: Mean. Blue: Minimum')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
# ax1.set_ylim(0,10)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# plt.legend()

```



Correlations between the temperatures

In [58]:

```

"""highly correlated at lag = 0"""

from scipy.stats import pearsonr

```



```
print(pearsonr(temperatures_yearly_first_diff['LandAverageTemperature'], temperatures_yearly_first_diff['LandMaxT
# (0.7041473540151104, 5.074132061908521e-26)

print(pearsonr(temperatures_yearly_first_diff['LandAverageTemperature'], temperatures_yearly_first_diff['LandMinT
# (0.8674877678792309, 2.675622965207922e-51)

print(pearsonr(temperatures_yearly_first_diff['LandMaxTemperature'], temperatures_yearly_first_diff['LandMinTempe
# (0.7165142434197547, 2.7704818163204947e-27)

(0.7041473540151104, 5.074132061908521e-26)
(0.8674877678792309, 2.675622965207922e-51)
(0.7165142434197547, 2.7704818163204947e-27)
```

So it looks like the three time series behave very similarly. For our purpose, we just study the 'LandAverageTemperature'. Moreover, we use data starting from 1850 from now on, due to the smaller uncertainty.

Augmented Dickey-Fuller test for non-stationarity

Non-stationarity is the null hypothesis; p-value > 0.05 infers non-stationary.

```
In [47]: from statsmodels.tsa.stattools import adfuller
```

```
In [48]: """# p-values > 0.05: non-stationary"""
print(adfuller(temperatures_1850_all_yearly['LandAverageTemperature'])[1])
print(adfuller(temperatures_1850_all_yearly['LandMaxTemperature'])[1])
print(adfuller(temperatures_1850_all_yearly['LandMinTemperature'])[1])

0.9505464412068493
0.8231808221341292
0.9814993947787101
```

```
In [9]: """# first differences"""
temperatures_yearly_first_diff = temperatures_1850_all_yearly[['LandAverageTemperature', 'LandMaxTemperature', 'L
# drop year 1850 (first row; NaN after first difference)
```

```
In [51]: temperatures_yearly_first_diff.head()
```

```
Out[51]:
```

	LandAverageTemperature	LandMaxTemperature	LandMinTemperature
dt			
1851-12-31	0.277917	-0.395667	0.239583
1852-12-31	-0.078417	0.316333	0.133083
1853-12-31	-0.058333	0.489250	-0.444500
1854-12-31	0.168667	0.090833	-0.130333
1855-12-31	-0.099750	-0.484250	-0.059333


```
In [50]: """# p-values for first differences < 0.05: stationary"""
print(adfuller(temperatures_yearly_first_diff['LandAverageTemperature'])[1])
print(adfuller(temperatures_yearly_first_diff['LandMaxTemperature'])[1])
print(adfuller(temperatures_yearly_first_diff['LandMinTemperature'])[1])

2.625626865464363e-25
6.030368572920849e-24
1.3653286686687314e-08
```

```
In [10]: """# plot first difference of yearly-averaged LandAverageTemperature and error bar 'LandAverageTemperatureUncerta

plt.figure(dpi=100)
plt.errorbar(temperatures_yearly_first_diff.index, 'LandAverageTemperature', data=temperatures_yearly_first_diff,
# plt.errorbar(temperatures_yearly_first_diff.index, 'LandMaxTemperature', data=temperatures_yearly_first_diff, c
# plt.errorbar(temperatures_yearly_first_diff.index, 'LandMinTemperature', data=temperatures_yearly_first_diff, c
ax1 = plt.gca()
# ax1.set_ylim(0,10)
```


```
plt.title('First Difference of Land Temperature Averaged by Year')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# ax1.get_legend().remove()
```



Auto-correlation function (ACF) and partial auto-correlation function (PACF)

```
In [54]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [56]: """# Auto-correlations"""
plt.figure(dpi=100)
plot_acf(temperatures_yearly_first_diff['LandAverageTemperature'], ax = plt.gca());
ax1 = plt.gca()
ax1.set_xlabel('lags/year')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# plt.savefig(r'plots/LandAverageTemperature_year_first_diff_ACF', dpi=150)
```




```
In [57]:
```

```

"""# Partial auto-correlations"""
plt.figure(dpi=100)
plot_pacf(temperatures_yearly_first_diff['LandAverageTemperature'], ax = plt.gca());
ax1 = plt.gca()
ax1.set_xlabel('lags/year')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# plt.savefig(r'plots/LandAverageTemperature_year_first_diff_PACF', dpi=150)

```



The decay in ACF, and the PACF, suggest AR(3).

Train-validation (initial) split

```

In [20]: # temperatures_yearly_first_diff.index[0]
# Timestamp('1851-12-31 00:00:00', freq='A-DEC')
# temperatures_yearly_first_diff.index[-1]
# Timestamp('2015-12-31 00:00:00', freq='A-DEC')

"""split at 1986; 1851-1985 is training set, 1986-2015 is validation set"""
split_time = pd.Timestamp(1986,6,30)
# temperatures_yearly_first_diff.index[136]
# Timestamp('1986-12-31 00:00:00', freq='A-DEC')
# 136

temperatures_1850_all_yearly_train = temperatures_1850_all_yearly[temperatures_1850_all_yearly.index < split_time]
temperatures_yearly_first_diff_train = temperatures_yearly_first_diff[temperatures_yearly_first_diff.index < split_time]

temperatures_1850_all_yearly_valid = temperatures_1850_all_yearly[temperatures_1850_all_yearly.index > split_time]
temperatures_yearly_first_diff_valid = temperatures_yearly_first_diff[temperatures_yearly_first_diff.index > split_time]

```

Naive forecast (lag 1)

```

In [22]: temperature_yearly_naive = [np.nan] + temperatures_1850_all_yearly['LandAverageTemperature'].to_list()[1:]

In [16]: from sklearn.metrics import mean_absolute_error

In [23]: mean_absolute_error(temperatures_1850_all_yearly_valid['LandAverageTemperature'], pd.Series(temperature_yearly_naive[1:]))

Out[23]: 0.18930555555555555

```

Auto-regressive (AR) models

```
In [60]: from sklearn.metrics import mean_absolute_error
# ARIMA models
from statsmodels.tsa.arima_model import ARIMA
```

```
In [61]: """fit model of the initial fixed partition"""
model1 = ARIMA(temperatures_yearly_first_diff_train['LandAverageTemperature'], order=(3,0,0))
model1_fit = model1.fit()

# summary of the model
print(model1_fit.summary())
```

```

                        ARMA Model Results
=====
Dep. Variable:      LandAverageTemperature    No. Observations:      135
Model:              ARMA(3, 0)               Log Likelihood        34.518
Method:              css-mle                  S.D. of innovations    0.187
Date:               Sun, 07 Nov 2021          AIC                   -59.035
Time:               00:01:57                  BIC                   -44.509
Sample:             12-31-1851                HQIC                  -53.132
                  - 12-31-1985
=====
                        coef      std err      z      P>|z|      [0.025      0.975]
-----
const                0.0053      0.007      0.749      0.454      -0.009      0.019
ar.L1.LandAverageTemperature -0.5168      0.079     -6.577      0.000      -0.671     -0.363
ar.L2.LandAverageTemperature -0.3708      0.085     -4.382      0.000      -0.537     -0.205
ar.L3.LandAverageTemperature -0.4115      0.079     -5.217      0.000      -0.566     -0.257
Roots
=====
                        Real      Imaginary      Modulus      Frequency
-----
AR.1                -1.3303      -0.0000j      1.3303      -0.5000
AR.2                 0.2146      -1.3345j      1.3516      -0.2246
AR.3                 0.2146      +1.3345j      1.3516      0.2246
=====
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\arima_model.py:472: FutureWarning: statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)

warnings.warn(ARIMA_DEPRECATION_WARN, FutureWarning)
```

Indeed, p-values are small (zero) at lag = 1, 2, 3.

[Fixed window] predictions/forecast

```
In [63]: """# in-sample (training set) prediction of first difference """
pred_diff_train1 = model1_fit.predict(start=temperatures_yearly_first_diff_train.index[0], end=temperatures_yearly_first_diff_train.index[-1])

# add back the trend
predictions_train1 = temperatures_1850_all_yearly_train['LandAverageTemperature'].shift() + pred_diff_train1
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:132: FutureWarning: The 'freq' argument in Timestamp is deprecated and will be removed in a future version.
date_key = Timestamp(key, freq=base_index.freq)

```
In [64]: """# mean_absolute_error of the training set """
mean_absolute_error(temperatures_1850_all_yearly_train['LandAverageTemperature'].iloc[1:], predictions_train1.iloc[1:])
```

```
Out[64]: 0.14860975236678103
```

[fixed window], out-of-sample forecast of first difference

```
In [65]: """ [fixed window], out-of-sample forecast of first difference """


forecast_diff1 = model1_fit.forecast(len(temperatures_yearly_first_diff_valid))[0]


fixed_wind_forecast1 = []
for i in range(len(temperatures_yearly_first_diff_valid)):
    if i == 0:
        fixed_wind_forecast1.append(temperatures_1850_all_yearly_train['LandAverageTemperature'][-1] + forecast_diff1)
    else:
        fixed_wind_forecast1.append(temperatures_1850_all_yearly_valid['LandAverageTemperature'][i-1] + forecast_diff1)

fixed_wind_forecast1 = pd.Series(fixed_wind_forecast1, index = temperatures_1850_all_yearly_valid.index)
```

```
In [76]: plt.figure(dpi=100)
temperatures_yearly_first_diff['LandAverageTemperature'].plot(ax=plt.gca())
pd.concat([pred_diff_train1, pd.Series(forecast_diff1, index = temperatures_yearly_first_diff_valid.index)], axis=1)
ax1 = plt.gca()
ax1.get_lines()[0].set_color("green")
ax1.get_lines()[1].set_color("tomato")
plt.title('First Difference of Land Temperature Averaged by Year: \nTrue values (green) VS fixed-partition model prediction (red)')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# plt.savefig(r'plots/LandAverageTemperature_year_AR3_fixed_partition_first_diff', dpi=150)

plt.figure(dpi=100)
temperatures_1850_all_yearly['LandAverageTemperature'].plot(ax=plt.gca())
pd.concat([predictions_train1, fixed_wind_forecast1], axis = 0).plot(ax=plt.gca())
ax1 = plt.gca()
ax1.get_lines()[0].set_color("green")
ax1.get_lines()[1].set_color("tomato")
plt.title('Land Temperature Averaged by Year: \nTrue values (green) VS fixed-partition model prediction (red)')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# ax1.get_legend().get_texts()[0].set_text('AR(3) fixed partition forecast')
# ax1.get_legend().get_texts()[1].set_text('True values')
# plt.savefig(r'plots/LandAverageTemperature_year_AR3_fixed_partition', dpi=150)
```





Validation MAE

```
In [67]: print(mean_absolute_error(temperatures_yearly_first_diff_valid['LandAverageTemperature'], pd.Series(forecast_diff
print(mean_absolute_error(temperatures_1850_all_yearly_valid['LandAverageTemperature'], fixed_wind_forecast1))

0.18836564357881894
0.1883656435788189
```

The forecast first difference of the fixed-partitioning model is basically zero; this reduces to a naive forecast by the value of the previous time-step.

Rolling forecast

Walk-forward validation

```
In [ ]: size = 136
# temperatures_1850_all_yearly.index[136]
# Out[159]: Timestamp('1986-12-31 00:00:00', freq='A-DEC')
history1 = temperatures_yearly_first_diff_train['LandAverageTemperature'].to_list()
rolling_forecast_first_diff1 = []

for t in range(len(temperatures_1850_all_yearly_valid)):
    model = ARIMA(history1, order=(3,0,0))
    model_fit = model.fit()
    output = model_fit.forecast() # outputs an array of 1 element
    yhat = output[0][0] # the forecast one time step ahead
    rolling_forecast_first_diff1.append(yhat)
    obs = temperatures_yearly_first_diff_valid['LandAverageTemperature'].iloc[t]
    history1.append(obs) # append value of current time step for next model
```

[rolling window], out-of-sample forecast from first difference

```
In [78]: """ [rolling window], out-of-sample forecast of first difference """

rolling_forecast1 = []
for i in range(len(temperatures_yearly_first_diff_valid)):
    if i == 0:
        rolling_forecast1.append(temperatures_1850_all_yearly_train['LandAverageTemperature'][-1] + rolling_forec
    else:
        rolling_forecast1.append(temperatures_1850_all_yearly_valid['LandAverageTemperature'][i-1] + rolling_forec

rolling_forecast1 = pd.Series(rolling_forecast1, index = temperatures_1850_all_yearly_valid.index)
```



```
In [80]: plt.figure(dpi=100)
```

```

temperatures_yearly_first_diff['LandAverageTemperature'].plot(ax=plt.gca())
pd.concat([pred_diff_train1, pd.Series(rolling_forecast_first_diff1, index = temperatures_yearly_first_diff_valid
ax1 = plt.gca()
ax1.get_lines()[0].set_color("green")
ax1.get_lines()[1].set_color("tomato")
plt.title('First Difference of Land Temperature Averaged by Year: \nTrue values (green) VS rolling-window model p
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
# plt.savefig(r'plots/LandAverageTemperature_year_AR3_rolling_forecast_first_diff', dpi=150)

plt.figure(dpi=100)
temperatures_1850_all_yearly['LandAverageTemperature'].plot(ax=plt.gca())
pd.concat([predictions_train1, rolling_forecast1], axis = 0).plot(ax=plt.gca())
ax1 = plt.gca()
ax1.get_lines()[0].set_color("green")
ax1.get_lines()[1].set_color("tomato")
plt.title('Land Temperature Averaged by Year: \nTrue values (green) VS rolling-window model prediction (red)')
ax1.set_xlabel('Year')
ax1.set_ylabel('Celsius')
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)

```



Validation MAE

```
In [79]: print(mean_absolute_error(temperatures_yearly_first_diff_valid['LandAverageTemperature'], pd.Series(rolling_forec
)
0.15548113397624228
```

Smaller error indeed, compared to fixed-partition forecast.

Summary of Results

Validation (data from 1986 onward) MAE:
Naive forecast (lag 1): 0.18930555555555575
AR(3) fixed partitioning: 0.1883656435788189
AR(3) rolling forecast: 0.15548113397624228

```
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
In [ ]:
```